

RAC

How do you list nodes of a cluster?

```
$ olsnodes
```

How do you do a health check of Oracle clusterware?

```
$ crsctl check crs  
$ crs_stat -t
```

What does a voting disk contain?

- * Node membership
- * Heartbeats

How do you find the voting disk location?

```
crsctl query css votedisk
```

How do you check what instances are running?

```
$ srvctl status database -d <db>
```

How do you check which listeners are running?

```
$ srvctl status nodeapps -n <node>  
$ crs_stat -t  
$ lsnrctl status      (you need to run this on every node)
```

What IP addresses do you have in a 10g RAC system?

- * Public IP (physical IP address)
- * Private IP (for interconnect)
- * Virtual IP

New in 10g. This is for connections. The virtual IP addresses are used to avoid timeouts when a node is down. The virtual IP is marked as "down" and connections move on to the next node immediately (no waits for timeouts).

What are Virtual IP addresses (VIPs) used by in RAC?

VIPs are related to a NIC card for failover and automatically managed by CRS .

When installing 10g RAC, at least 2 network interfaces are required for each node:
One for public and one for private Interface.

How do you identify the interface name and associated IP address for every network adapter?

```
$ ifconfig -a
```

What is the past image of a block?

RAC cache fusion:

Past images are older copies of a buffer that have not yet been written to disk.

If a lock is acquired and already has dirty buffers on a remote instance then it takes on a global lock role and any remote instances that contain these dirty buffers will have what is called a 'past image' of the buffer.

What files are shared (or not shared) in a RAC system?

Common/shared files (shared)

- data files
- control files

Instance specific files (not shared)

- redo log files
 - Put on shared storage to help recovery operations.
- archive log files
 - Put on shared or local storage (better on shared).
- undo tablespace
- alert.log and trace files

Special cases

- spfile

RAC systems can have a single shared spfile when using a symbolic link in \$ORACLE_HOME/dbs pointing to the shared file.

What is Cache Fusion?

Cache Fusion provides cache-to-cache transfer of data blocks between instances in a cluster.

What happens if you create a datafile on a file system that is only visible to one node (ocfs)?

Most things will still work (e.g. transactions that don't need to access the invisible data file). You will get an error for transactions that try to access the invisible data file.

What are the CRS components?

- * Voting disk to record node membership information and heartbeats.
- * Oracle Cluster Registry (OCR) to record cluster configuration information

How do you list the private interconnect?

```
olsnodes -p -i           # shows private interconnect and virtual host name
```

How do you check if the CRS processes are running?

```
$ ps -ef |grep crsd          # cluster ready services  
$ ps -ef |grep cssd          # cluster synchronisation service  
$ ps -ef |grep evmd          # event manager
```

How do you export, import, backup and restore OCR?

```
ocrconfig -showbackup
```

How do you start/stop an Oracle RAC system?

Starting the Oracle RAC 10g Environment

1. Start the Oracle clusterware as root
2. Start the node applications (Virtual IP, GSD, TNS Listener, and ONS).
3. Bring up the ASM instance.
4. Start up the Oracle instance (and related services).
5. Start the Enterprise Manager Database console.

```
$ crsctl start crs           # as root  
$ srvctl start nodeapps -n node  
$ srvctl start asm -n node  
$ srvctl start instance -d db_name -i instance  
$ emctl start dbconsole
```

Stopping the Oracle RAC 10g Environment

Run the stop commands in reverse order of the start commands.

How do you add/remove a database or instance to CRS?

```
srvctl add database -d <db> -o <oracle_home>  
srvctl remove database -d <db> -f  
  
srvctl add instance -d <db> -i <sid> -n <node>  
srvctl remove instance -d <db> -i <sid>
```

How do you start/stop listener and database in a RAC system?

```
srvctl start listener -n <node>
srvctl stop listener -n <node>

srvctl start database -d <db>
srvctl stop database -d <db>
```

What are RAC specific tuning issues?

- Block pinning between instances.
- Single threaded batch job runs only on one instance.
- You need to know on which instance a problem SQL is running in order to tune it.
- Check log and trace files on all nodes.

What are RAC specific administration issues?

- Need to check alert.log file for each instance.
- Use srvctl instead of sqlplus to start/stop the instance/database
- Parameters can be set globally (*.param) or instance specific (instance.param)
- Cluster related issues, e.g. node eviction
- Scan listener

How do you upgrade a RAC database?

1. Upgrade Clusterware
2. Upgrade ASM
3. Upgrade the database

Startup the database in single instance mode for the upgrade.

Do a rolling upgrade of the databases.

What are the CRS background Processes?

oprocda

Process monitor for the cluster.

evmd

Event manager daemon that starts the racgevt process to manage callouts.

ocssd

Manages cluster node membership and runs as oracle user; failure of this process results in cluster restart.

crsd

Performs high availability recovery and management operations such as maintaining the OCR. Also manages application resources and runs as root user and restarts automatically upon failure.

Table of Contents

How do you list nodes of a cluster?.....	1
How do you do a health check of Oracle clusterware?.....	1
What does a voting disk contain?.....	1
How do you find the voting disk location?.....	1
How do you check what instances are running?.....	1
How do you check which listeners are running?.....	1
What IP addresses do you have in a 10g RAC system?.....	1
What are Virtual IP addresses (VIPs) used by in RAC?.....	1
How do you identify the interface name and associated IP address for every network adapter?.....	2
What is the past image of a block?.....	2
What files are shared (or not shared) in a RAC system?.....	2
What is Cache Fusion?.....	2
What happens if you create a datafile on a file system that is only visible to one node (ocfs)?.....	2
What are the CRS components?.....	3
How to you list the private interconnect?.....	3
How do you check if the CRS processes are running?.....	3
How to you export, import, backup and restore OCR?.....	3
How to you start/stop an Oracle RAC system?.....	3
How do you add/remove a database or instance to CRS?.....	3
How do you start/stop listener and database in a RAC system?.....	4
What are RAC specific tuning issues?.....	4
What are RAC specific administration issues?.....	4
How do you upgrade a RAC database?.....	4
What are the CRS background Processes?.....	4